

Case Study 6

Hollie Gardner

Assignment Instructions

Build a dense neural network to accurately detect the particle. The goal is to maximize your accuracy. Include a discussion of how you know your model has finished training as well as what design decisions you made while building the network.

```
In [ ]: # libraries from videos
import pandas as pd
import numpy as np
import os

#model prep
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

# modeling
import tensorflow as tf

# example of training a final classification model
from keras.models import Sequential
from keras.layers import Dense

#evaluation metrics
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix

#data prep and viz
import seaborn as sns
import matplotlib.pyplot as plt
```

Initial EDA

```
In [ ]: #read in the data
os.chdir('/Users/holliegardner/Desktop/qtw6/')
df = pd.read_csv('all_train.csv')
```

```
In [ ]: df.shape
```

```
Out[ ]: (7000000, 29)
```

Dataset Overview

- 7 million observations
- 1 dependent variable
- 28 features

In []:

`df.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7000000 entries, 0 to 6999999
Data columns (total 29 columns):
#   Column      Dtype
---  -
0   # label     float64
1   f0          float64
2   f1          float64
3   f2          float64
4   f3          float64
5   f4          float64
6   f5          float64
7   f6          float64
8   f7          float64
9   f8          float64
10  f9          float64
11  f10         float64
12  f11         float64
13  f12         float64
14  f13         float64
15  f14         float64
16  f15         float64
17  f16         float64
18  f17         float64
19  f18         float64
20  f19         float64
21  f20         float64
22  f21         float64
23  f22         float64
24  f23         float64
25  f24         float64
26  f25         float64
27  f26         float64
28  mass        float64
dtypes: float64(29)
memory usage: 1.5 GB

```

All features came in as float even though the dependent variable should be integer type.
Changing in the next cell.

In []:

```

#changing target from float to integer
df['# label'] = df['# label'].astype(int)
print(df['# label'])

```

```

0          1
1          1
2          0
3          1
4          0
..
6999995    0
6999996    0
6999997    1
6999998    1
6999999    1
Name: # label, Length: 7000000, dtype: int64

```

In []:

```

# counts of each class
df['# label'].value_counts()

```

```
Out[ ]: 1    3500879
        0    3499121
        Name: # label, dtype: int64
```

```
In [ ]: #separating target values from dataset
X = df.drop(['# label'], axis=1)
y = df['# label']
```

```
In [ ]: #test train split 60/40
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=.2, random_state=143)

#standardizing data
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
In [ ]: #sequential model - sequential class
model = tf.keras.Sequential()

#use the keras api call for input layer.
model.add(tf.keras.Input(shape=(28,)))

#hidden layer
model.add(tf.keras.layers.Dense(50, activation='relu'))
model.add(tf.keras.layers.Dense(50, activation='relu'))

#add the output layer, match final layer with dimension.
model.add(tf.keras.layers.Dense(1, activation='sigmoid'))
```

Resources for selecting parameters

<https://medium.com/geekculture/a-2021-guide-to-improving-cnns-optimizers-adam-vs-sgd-495848ac6008>

```
In [ ]: # compiling
model.compile(
    optimizer='adam',
    loss=tf.keras.losses.BinaryCrossentropy(),
    metrics=['accuracy']
)
```

```
In [ ]: #fit the model with -- data, target, # of epochs
model.fit(X_train_scaled, y_train, epochs=10, batch_size=100000)
```

```
Epoch 1/10
56/56 [=====] - 10s 133ms/step - loss: 0.5370 - accurac
y: 0.7196
Epoch 2/10
56/56 [=====] - 7s 131ms/step - loss: 0.3799 - accurac
y: 0.8249
Epoch 3/10
56/56 [=====] - 7s 132ms/step - loss: 0.3470 - accurac
```

```

y: 0.8396
Epoch 4/10
56/56 [=====] - 7s 132ms/step - loss: 0.3314 - accurac
y: 0.8463
Epoch 5/10
56/56 [=====] - 7s 128ms/step - loss: 0.3222 - accurac
y: 0.8505
Epoch 6/10
56/56 [=====] - 7s 133ms/step - loss: 0.3148 - accurac
y: 0.8539
Epoch 7/10
56/56 [=====] - 7s 128ms/step - loss: 0.3081 - accurac
y: 0.8568
Epoch 8/10
56/56 [=====] - 8s 134ms/step - loss: 0.3029 - accurac
y: 0.8590
Epoch 9/10
56/56 [=====] - 8s 135ms/step - loss: 0.2990 - accurac
y: 0.8605
Epoch 10/10
56/56 [=====] - 7s 128ms/step - loss: 0.2960 - accurac
y: 0.8617

```

Out[]: <keras.callbacks.History at 0x1068dc760>

```

In [ ]: #creating predictions for test holdout.
y_pred = (model.predict(X_test_scaled) >= 0.50).astype(int)

```

```

43750/43750 [=====] - 68s 2ms/step

```

```

In [ ]: #quick visual check of predictions
print(y_pred)

```

```

[[1]
 [0]
 [0]
 ...
 [0]
 [1]
 [1]]

```

```

In [ ]: # evaluation metrics
print(classification_report(y_test, y_pred))

```

	precision	recall	f1-score	support
0	0.89	0.83	0.86	700525
1	0.84	0.89	0.87	699475
accuracy			0.86	1400000
macro avg	0.86	0.86	0.86	1400000
weighted avg	0.86	0.86	0.86	1400000

```

In [ ]: #Generate a visual confusion matrix
cf_matrix = confusion_matrix(y_test, y_pred)
sns.set(font_scale=1.4)
ax = sns.heatmap(cf_matrix, linewidths=.5, fmt='.5g', annot=True, cmap="coolwarm")
ax.set_title('Model 1: Confusion Matrix\n\n')
ax.set_xlabel('\nPredicted Values')
ax.set_ylabel('Actual Values ')

```

```
ax.xaxis.set_ticklabels(['Noise', 'Signal'])  
ax.yaxis.set_ticklabels(['Noise', 'Signal'])  
plt.show()
```

